

Model Selection - Best Subset

This is a pipeline of model selection using Best Subset method. It includes the *R* codes for *k*-folder cross validation using the prostate dataset.

Step 1 - Data Loading

```
source("subset_selection.R")

X = load_data("Data/prostate.data")

train_data = X$train_data
test_data = X$test_data
n = X$num_of_train_sample
p = X$num_of_predictors
```

Step 2 - Predicting

Scaling

separate the predictors and the response

```
y_train = train_data[,c(response_name)]
train_data = train_data[,c(1:p)]
y_test = test_data[,c(response_name)]
test_data = test_data[,c(1:p)]
```

```
head(train_data, n=5)
```

##	lcavol	lweight	age	lbph	svi	lcp	gleason	pgg45	lpsa
## 1	-0.5798185	2.769459	50	-1.386294	0	-1.386294	6	0	-0.4307829
## 2	-0.9942523	3.319626	58	-1.386294	0	-1.386294	6	0	-0.1625189
## 3	-0.5108256	2.691243	74	-1.386294	0	-1.386294	7	20	-0.1625189
## 4	-1.2039728	3.282789	58	-1.386294	0	-1.386294	6	0	-0.1625189
## 5	0.7514161	3.432373	62	-1.386294	0	-1.386294	6	0	0.3715636

scale:

$$x = \frac{(x - \text{mean}(x))}{\text{std}(x)}$$

1. scale training data (without labels)
2. demean training labels & combine the training dataset
3. manually scale testing data
4. demean testing labels (using the mean of the training label) & combine the testing dataset

```
scaling <- function(train_data, test_data, p){
```

```
  y_train = train_data[,c(response_name)]
  train_data = train_data[,c(1:p)]
```

```
  y_test = test_data[,c(response_name)]
```

```

test_data=test_data[,c(1:p)]

train_data      = scale(train_data)
means           = attr(train_data, "scaled:center")
stds            = attr(train_data, "scaled:scale")

y_train_mean    = mean(y_train)
y_train         = y_train - y_train_mean
train_data      = data.frame(cbind(train_data, y_train))
names(train_data)[p+1] = response_name

test_data       = t(apply(test_data, 1, '-', means))
test_data       = t(apply(test_data, 1, '/', stds))
test_data       = data.frame(cbind(test_data, y_test - y_train_mean))
names(test_data)[p+1] = response_name

return(list(train_data=train_data, test_data=test_data,
            y_train_mean=y_train_mean, y_test=y_test))
}

datasets      = scaling(train_data, test_data, p)
train_data    = datasets$train_data
test_data     = datasets$test_data
y_test        = datasets$y_test
y_train_mean  = datasets$y_train_mean

head(y_test, n=5)

## [1] 0.7654678 1.0473190 1.0473190 1.3987169 1.6582281

```

Predicting

```

predictor_formula = "lpsa ~ lcavol + lweight + age + lbph+ svi + lcp + gleason"
mk = lm(formula = predictor_formula, data = train_data)
y_hat = predict(mk, newdata = test_data, interval="prediction")[,1]
y_hat = y_hat + y_train_mean

```

Create a predicting function that will be used by Cross Validation

```

predicting <- function(predictor_formula, train_data, test_data, p){

  datasets      = scaling(train_data, test_data, p)
  train_data    = datasets$train_data
  test_data     = datasets$test_data
  y_test        = datasets$y_test
  y_train_mean  = datasets$y_train_mean

  mk = lm(formula = predictor_formula, data = train_data)
  pdt = predict(mk, newdata = test_data, interval="prediction")[,1]
  pdt = pdt + y_train_mean

  return(list(pdt = pdt, y_mean = y_train_mean, y_test = y_test,
            n_test = length(y_test), coef=coef(mk)))
}

```

```
}
```

Step 3 - Cross Validation

LOOP (1:fold):

split the *data* into test and train datasets

obtain y_{hat} and y for given *predicting formula*

calculate *CV Squared Error*

```
cross_validation <- function(formula, data, fold=10){  
  n = dim(data)[1]  
  p = dim(data)[2] - 1  
  
  fold_size = round(n/fold)  
  
  for (i in 1:fold) {  
    test_index      = fold_size * (i-1) + 1:fold_size  
    test_index      = intersect(test_index, 1:n)  
    train_index     = setdiff(1:n, test_index)  
  
    cv_test         = data[test_index,]  
    cv_train        = data[train_index,]  
  
    results = predicting(formula, cv_train, cv_test, p)  
    pdt      = results$pdt  
    y_test   = results$y_test  
  
    y_hat    <- if (i==1) pdt else c(y_hat, pdt)  
    y        <- if (i==1) y_test else c(y, y_test)  
  }  
  
  cv_se  = (y - y_hat)^2  
  cv_mean = mean(cv_se)  
  cv_sd   = sqrt(var(cv_se)/length(y))  
  return(list(cv_squared_error = cv_se, cv_mean = cv_mean,  
             cv_sd= cv_sd, name="Mean Squared Error"))  
}
```

Step 4 - Search for Best Subset

Best Subset for a given number (k) of predictors

calculate *all possible subsets* given k

Loop (in possible subsets):

build a predicting formula

Cross Validation (formula, data, fold=10)

record *squared error*, *formula* and *model features*.

```
best_subset <- function(num_of_predictors, total_predictors, data, cv_fold=10){
  p = total_predictors
  stopifnot((num_of_predictors>=0) && (num_of_predictors<=p))

  best_error = 1e5

  if (num_of_predictors == 0){
    form      = paste(response_name, "~+1")
    result    = cross_validation(form, data, cv_fold)

    best_error = result$cv_squared_error
    best_features = NULL
    best_formula = form
  }
  else{
    all_possible_subsets = combn(p, num_of_predictors)
    num_of_subsets      = dim(all_possible_subsets)[2]

    for(i in 1:num_of_subsets){
      print(sprintf("i=%5d; num of subsets=%5d; percent finish=%5.6f", i,
                    num_of_subsets, i/num_of_subsets))

      feature_indices = all_possible_subsets[, i]
      features        = as.vector(names(data))[feature_indices]

      form            = paste(response_name, " ~ ")
      form            = paste(form, paste(features, collapse = '+'))
      result          = cross_validation(form, data, cv_fold)
      squared_error   = result$cv_squared_error

      if( mean(squared_error) < mean(best_error) ){
        best_error    = squared_error
        best_features  = feature_indices
        best_formula   = form
      }
    }
  }
  return(list(best_error, best_features, best_formula))
}
```

Find the best subset for all possible numbers of predictors

LOOP (k in 0:p):

Best_Subset(k, p, train data, cv fold=10)

One standard error rule

```

one_standard_error_rule <- function(complexity_parameter, cv_results){

  n = dim(cv_results)[1]

  means = apply(cv_results, 2, mean)
  stds = sqrt(apply(cv_results, 2, var) / n)

  # find the smallest espe:
  min_epe_index = which.min(means)

  # compute the confidence interval width around this point:
  ciw = stds[min_epe_index]

  threshold = means[min_epe_index] + 0.5*ciw

  complexity_index = which.min(abs(means - threshold))

  complexity_value = complexity_parameter[complexity_index]

  return(list(complexity = complexity_value,
              index = complexity_index,
              means = means,
              stds = stds))
}

```

Step 5 - Predict using the best model

```

# OSE = one_standard_error_rule(complexity_parameter, cv_results)
# complexity_parameter_idx = OSE$index
#
# predictor_formula = best_formula[complexity_parameter_idx]
#
# result = predicting(predictor_formula, train_data, test_data, 8)
# y = result$y_test
# pdt = result$pdt
# n = result$n_test
#
# mse = mean((y - pdt)^2)
# sErr = sqrt( var((y - pdt)^2) / n)

```